

● llama_v3_2_3b_instruct_prompt_1_of_3

Information

Name

llama_v3_2_3b_instruct_prompt_1_of_3

Target Device

Snapdragon X Elite CRD
Windows 11
Snapdragon® X Elite | SC8380XP

Creator

leoliu5550@gmail.com

Target Model

job_jgkqolv2g_linked_bin

Input Specs

input_ids : int32[1, 128]

Submission / Completion Time

2025/8/10 上午14:3:24
2025/8/10 上午1:51:30

Options

--max_profiler_iterations 50 --
qnn_options
context_enable_graphs=prompt_ar128_cl4096_1_

Versions

QAIRT : v2.34.2.25052816411119506
QNN Backend API : 5.34.0
QNN Core API : 2.25.0
Windows : Windows 11 (26I00)
Build ID : APSS.WP.HA.1.0-08200-
SC8380XRELSFNWZA-3
AI Hub : aihub-2025.0725.0

Inference Metrics

Minimum Inference Time

0.2 ms

Estimated Peak Memory Usage

0.0 MB

Compute Units

NPU

4

Detailed Metrics

Stage	Time	Memory
Compilation	0.0 ms	0.0 MB
First App Load	1048.0 ms	12 MB
Subsequent App Load	1004.4 ms	12 MB
Inference	(min) 0.2 ms (median) 0.2 ms 	0.0 MB

• llama_v3_2_3b_instruct_prompt_2_of_3

Information

Name

llama_v3_2_3b_instruct_prompt_2_of_3

Target Device

Snapdragon X Elite CRD
Windows 11
Snapdragon® X Elite | SC8380XP

Creator

leoliu5550@gmail.com

Target Model

job_jgj2o0zx5_linked_binmq3zv5j6m

Input Specs

_model_model_embed_tokens_gather_output_0 :

uint16[1, 128, 3072]

attention_mask : uint16[1, 1, 128, 4096]

position_ids_cos : uint16[1, 1, 128, 64]

position_ids_sin : uint16[1, 1, 128, 64]

past_key_0_in : uint8[8, 1, 128, 3968]

past_value_0_in : uint8[8, 1, 3968, 128]

past_key_1_in : uint8[8, 1, 128, 3968]

past_value_1_in : uint8[8, 1, 3968, 128]

Submission / Completion Time

2025/8/10 上午14:3:25
2025/8/10 上午1:52:16

Inference Metrics

Minimum Inference Time ⓘ

62.2 ms

Estimated Peak Memory Usage ⓘ

111 MB

Compute Units ⓘ

NPU

7620

Detailed Metrics

Stage	Time	Memory
Compilation ⓘ	0.0 ms	0.0 MB
First App Load ⓘ	1607.0 ms	20 MB
Subsequent App Load ⓘ	649.4 ms	20 MB
Inference ⓘ	(min) 62.2 ms (median) 63.2 ms	111 MB

• llama_v3_2_3b_instruct_prompt_3_of_3

Profile Job Results

Jobs > jgz8xxop

Results Ready

SHARE

Information

Name

llama_v3_2_3b_instruct_prompt_3_of_3

Target Device

Snapdragon X Elite CRD
Windows 11
Snapdragon® X Elite | SC8380XP

Creator

leoliu5550@gmail.com

Target Model

job.jgz8xjqp_linked_bin

mnj2pe1lm

Input Specs

`_model_model_layers_13_Add_1_output_0 :`
uint16[1, 128, 3072]
attention_mask : uint16[1, 1, 128, 4096]
padding : int64[1, 1, 128, 4096]

Submission / Completion Time

2025/8/10 上午2:16:25
2025/8/10 上午2:26:46

Inference Metrics

Minimum Inference Time ⓘ

74.3 ms

Estimated Peak Memory Usage ⓘ

111 MB

Compute Units ⓘ

NPU

7627

Detailed Metrics

Stage	Time	Memory
Compilation ⓘ	0.0 ms	0.0 MB
First App Load ⓘ	1922.2 ms	52 MB
Subsequent App Load ⓘ	1191.8 ms	52 MB
Inference ⓘ	(min) 74.3 ms (median) 75.5 ms	111 MB

● llama_v3_2_3b_instruct_token_1_of_3

Jobs > j5w38dd3p

Results Ready

SHARE

Information

Name

llama_v3_2_3b_instruct_token_1_of_3

Target Device

Snapdragon X Elite CRD
Windows 11
Snapdragon® X Elite | SC8380XP

Creator

leoliu5550@gmail.com

Target Model

job.jgkqolv2g_linked_bin

Input Specs

input_ids: int32[1, 1]

Submission / Completion Time

2025/8/10 上午2:16:26
2025/8/10 上午2:24:45

Options

--max_profiler_iterations 50 --
qnn_options
context_enable_graphs=token_ar1_cl4096_1_of_

Versions

QAIRT: v2.34.2.250528164111_119506
QNN Backend API: 5.34.0
QNN Core API: 2.25.0
Windows: Windows 11 (26100)
Build ID: APSS.WP_HA.1.0-08200-
SC8380XRELSFNWZA-3
AI Hub: aihub-2025.07.25.0

Inference Metrics

Minimum Inference Time

0.1 ms

Estimated Peak Memory Usage

0.0 MB

Compute Units


NPU

4

Detailed Metrics

Stage	Time	Memory
Compilation	0.0 ms	0.0 MB
First App Load	1061.5 ms	12 MB
Subsequent App Load	947.2 ms	12 MB
Inference	(min) 0.1 ms (median) 0.2 ms 	0.0 MB

● llama_v3_2_3b_instruct_token_2_of_3



● llama_v3_2_3b_instruct_token_3_of_3

Jobs > jplw7dd8g

Results Ready

SHARE

Information

Name

llama_v3_2_3b_instruct_token_3_of_3

Target Device

Snapdragon X Elite CRD
Windows 11
Snapdragon® X Elite | SC8380XP

Creator

leoliu5550@gmail.com

Target Model

job.jgzj8xjqp.linked_bin

mng2petlm

Input Specs

_model_model_layers_13_Add_1_output_0 :

uint16[1, 1, 3072]

attention_mask : uint16[1, 1, 4096]

position_ids_cos : uint16[1, 1, 64]

Submission / Completion Time

2025/8/10 上午2:16:27
2025/8/10 上午2:38:16

Inference Metrics

Minimum Inference Time

43.6 ms

Estimated Peak Memory Usage

112 MB

Compute Units

NPU

7583

Detailed Metrics

Stage	Time	Memory
Compilation	0.0 ms	0.0 MB
First App Load	1923.4 ms	17 MB
Subsequent App Load	1201.3 ms	17 MB
Inference	(min) 43.6 ms (median) 45.3 ms	112 MB